

# House Price Prediction (Regression)

March 18, 2025

## 1 1. Import Libraries and Dataset

### 1.1 a. Import Libraries

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

from scipy.stats import norm
from scipy import stats

import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

### 1.2 b. Read the Dataset

Data source: <https://www.kaggle.com/competitions/house-prices-advanced-regression-techniques/data>

```
[6]: # Read the dataset:
df_train = pd.read_csv('house_train.csv')
df_test = pd.read_csv('house_test.csv')
```

```
[8]: df_train
```

```
[8]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	
...	...	...	...	...	...	...	...	...	
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg	
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg	
1457	1458	70	RL	66.0	9042	Pave	NaN	Reg	

1458	1459	20	RL	68.0	9717	Pave	NaN	Reg
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	\
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
...	...	...	...	...	...	...	...	...	
1455	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1456	Lvl	AllPub	...	0	NaN	MnPrv	NaN	0	
1457	Lvl	AllPub	...	0	NaN	GdPrv	Shed	2500	
1458	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1459	Lvl	AllPub	...	0	NaN	NaN	NaN	0	

	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	2	2008	WD	Normal	208500
1	5	2007	WD	Normal	181500
2	9	2008	WD	Normal	223500
3	2	2006	WD	Abnorml	140000
4	12	2008	WD	Normal	250000
...	...	...	...	...	...
1455	8	2007	WD	Normal	175000
1456	2	2010	WD	Normal	210000
1457	5	2010	WD	Normal	266500
1458	4	2010	WD	Normal	142125
1459	6	2008	WD	Normal	147500

[1460 rows x 81 columns]

```
[10]: df_test
```

```
[10]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1461	20	RH	80.0	11622	Pave	NaN	Reg	
1	1462	20	RL	81.0	14267	Pave	NaN	IR1	
2	1463	60	RL	74.0	13830	Pave	NaN	IR1	
3	1464	60	RL	78.0	9978	Pave	NaN	IR1	
4	1465	120	RL	43.0	5005	Pave	NaN	IR1	
...	...	...	...	...	...	...	...	...	
1454	2915	160	RM	21.0	1936	Pave	NaN	Reg	
1455	2916	160	RM	21.0	1894	Pave	NaN	Reg	
1456	2917	20	RL	160.0	20000	Pave	NaN	Reg	
1457	2918	85	RL	62.0	10441	Pave	NaN	Reg	
1458	2919	60	RL	74.0	9627	Pave	NaN	Reg	

LandContour	Utilities	...	ScreenPorch	PoolArea	PoolQC	Fence	\
-------------	-----------	-----	-------------	----------	--------	-------	---

0	Lvl	AllPub	...	120	0	NaN	MnPrv
1	Lvl	AllPub	...	0	0	NaN	NaN
2	Lvl	AllPub	...	0	0	NaN	MnPrv
3	Lvl	AllPub	...	0	0	NaN	NaN
4	HLS	AllPub	...	144	0	NaN	NaN
...	...	...	...	...	...	...	...
1454	Lvl	AllPub	...	0	0	NaN	NaN
1455	Lvl	AllPub	...	0	0	NaN	NaN
1456	Lvl	AllPub	...	0	0	NaN	NaN
1457	Lvl	AllPub	...	0	0	NaN	MnPrv
1458	Lvl	AllPub	...	0	0	NaN	NaN

	MiscFeature	MiscVal	MoSold	YrSold	SaleType	SaleCondition
0	NaN	0	6	2010	WD	Normal
1	Gar2	12500	6	2010	WD	Normal
2	NaN	0	3	2010	WD	Normal
3	NaN	0	6	2010	WD	Normal
4	NaN	0	1	2010	WD	Normal
...	...	...	...	...	...	...
1454	NaN	0	6	2006	WD	Normal
1455	NaN	0	4	2006	WD	Abnorml
1456	NaN	0	9	2006	WD	Abnorml
1457	Shed	700	7	2006	WD	Normal
1458	NaN	0	11	2006	WD	Normal

[1459 rows x 80 columns]

### 1.3 c. Selecting Important columns on the Dataframe

```
[13]: print(df_train.columns)
```

```
Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
      'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
      'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
      'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
      'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
      'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
      'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
      'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
      'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
      'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
      'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
      'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
      'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
      'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
      'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
      'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
```

```
    'SaleCondition', 'SalePrice'],
    dtype='object')
```

```
[15]: # Selecting the important columns from two DataFrames:
df_train = df_train[['LotArea', 'YearBuilt', 'OverallQual', 'OverallCond',
    ↳ 'RoofMatl', 'ExterQual', 'BedroomAbvGr', 'GrLivArea', 'MiscFeature',
    ↳ 'GarageArea', 'KitchenQual', 'SaleCondition', 'SalePrice']]
df_test = df_test[['LotArea', 'YearBuilt', 'OverallQual', 'OverallCond',
    ↳ 'RoofMatl', 'ExterQual', 'BedroomAbvGr', 'GrLivArea', 'MiscFeature',
    ↳ 'GarageArea', 'KitchenQual', 'SaleCondition']]
```

```
[17]: df_train
```

```
[17]:
```

	LotArea	YearBuilt	OverallQual	OverallCond	RoofMatl	ExterQual	\
0	8450	2003	7	5	CompShg	Gd	
1	9600	1976	6	8	CompShg	TA	
2	11250	2001	7	5	CompShg	Gd	
3	9550	1915	7	5	CompShg	TA	
4	14260	2000	8	5	CompShg	Gd	
...	...	...	...	...	...	...	
1455	7917	1999	6	5	CompShg	TA	
1456	13175	1978	6	6	CompShg	TA	
1457	9042	1941	7	9	CompShg	Ex	
1458	9717	1950	5	6	CompShg	TA	
1459	9937	1965	5	6	CompShg	Gd	

	BedroomAbvGr	GrLivArea	MiscFeature	GarageArea	KitchenQual	\
0	3	1710	NaN	548	Gd	
1	3	1262	NaN	460	TA	
2	3	1786	NaN	608	Gd	
3	3	1717	NaN	642	Gd	
4	4	2198	NaN	836	Gd	
...	...	...	...	...	...	
1455	3	1647	NaN	460	TA	
1456	3	2073	NaN	500	TA	
1457	4	2340	Shed	252	Gd	
1458	2	1078	NaN	240	Gd	
1459	3	1256	NaN	276	TA	

	SaleCondition	SalePrice
0	Normal	208500
1	Normal	181500
2	Normal	223500
3	Abnorml	140000
4	Normal	250000
...	...	...
1455	Normal	175000

1456	Normal	210000
1457	Normal	266500
1458	Normal	142125
1459	Normal	147500

[1460 rows x 13 columns]

## 2 2. Exploratory Data Analysis (EDA) for Train Set

### 2.1 a. Check Dataset Information

```
[21]: # Check data information:
df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   LotArea          1460 non-null   int64
1   YearBuilt        1460 non-null   int64
2   OverallQual      1460 non-null   int64
3   OverallCond      1460 non-null   int64
4   RoofMatl        1460 non-null   object
5   ExterQual        1460 non-null   object
6   BedroomAbvGr     1460 non-null   int64
7   GrLivArea        1460 non-null   int64
8   MiscFeature      54 non-null     object
9   GarageArea       1460 non-null   int64
10  KitchenQual      1460 non-null   object
11  SaleCondition    1460 non-null   object
12  SalePrice        1460 non-null   int64
dtypes: int64(8), object(5)
memory usage: 148.4+ KB
```

- There is missing value on MiscFeature

```
[24]: # Check data statistical summary:
df_train.describe()
```

```
[24]:
```


	LotArea	YearBuilt	OverallQual	OverallCond	BedroomAbvGr	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	10516.828082	1971.267808	6.099315	5.575342	2.866438	
std	9981.264932	30.202904	1.382997	1.112799	0.815778	
min	1300.000000	1872.000000	1.000000	1.000000	0.000000	
25%	7553.500000	1954.000000	5.000000	5.000000	2.000000	
50%	9478.500000	1973.000000	6.000000	5.000000	3.000000	

75%	11601.500000	2000.000000	7.000000	6.000000	3.000000
max	215245.000000	2010.000000	10.000000	9.000000	8.000000

	GrLivArea	GarageArea	SalePrice
count	1460.000000	1460.000000	1460.000000
mean	1515.463699	472.980137	180921.195890
std	525.480383	213.804841	79442.502883
min	334.000000	0.000000	34900.000000
25%	1129.500000	334.500000	129975.000000
50%	1464.000000	480.000000	163000.000000
75%	1776.750000	576.000000	214000.000000
max	5642.000000	1418.000000	755000.000000

## 2.2 b. Histogram of Target Variable (Sale Price)

```
[27]: # Show distribution of "SalePrice" column:
sns.histplot(df_train['SalePrice'])
plt.show()
```



- The sale prices is not normally distributed.
- The data appears to be skewed (positively/negatively) and may contain outliers.

- This suggests that the sale prices are not evenly distributed around the mean, which could impact statistical analyses that assume normality. So, preprocessing is needed before further analysis.

```
[29]: # Calculate skewness and kurtosis values:
# Skewness: a statistical measure used to assess how much the data distribution
↳ is skewed or asymmetric
# Kurtosis: a statistical measure used to assess how much the data distribution
↳ has long tails (outliers) and a sharp peak (peakedness) compared to a normal
↳ distribution (Gaussian distribution or bell-shaped distribution)

print (f"Skewness: {df_train['SalePrice'].skew()}")
print (f"Kurtosis: {df_train['SalePrice'].kurt()}")
```


Skewness: 1.8828757597682129

Kurtosis: 6.536281860064529

- The skewness value of 1.88 suggests that the data is positively skewed (right-skewed). This means that most of the sale prices are clustered towards the lower end, with a few higher-priced outliers.
- In other words, there are more lower-priced properties, and fewer high-priced properties that stretch the distribution to the right.
- The kurtosis value of 6.54 is significantly greater than 3 (which represents a normal distribution's kurtosis). This indicates that the distribution has heavy tails and a sharper peak, which means there are outliers or extreme values present in your data.

## 2.3 c. Scatter Plot of GrLivArea vs Target Variable (Sale Price)

```
[34]: df_train.plot.scatter(x='GrLivArea', y='SalePrice')
plt.show()
```



- The larger the area, the price tends to be the higher .
- There is an anomaly where large areas have low prices.

[37]: *# Check the data anomaly:*

```
df_train[(df_train['GrLivArea'] > 4000) & (df_train['SalePrice'] < 200000) |
          ((df_train['GrLivArea'] > 4000) & (df_train['SalePrice'] > 700000))]
```

```
[37]:
```

	LotArea	YearBuilt	OverallQual	OverallCond	RoofMatl	ExterQual	\
523	40094	2007	10	5	CompShg	Ex	
691	21535	1994	10	6	WdShngl	Ex	
1182	15623	1996	10	5	CompShg	Gd	
1298	63887	2008	10	5	ClyTile	Ex	

	BedroomAbvGr	GrLivArea	MiscFeature	GarageArea	KitchenQual	\
523	3	4676	NaN	884	Ex	
691	4	4316	NaN	832	Ex	
1182	4	4476	NaN	813	Ex	
1298	3	5642	NaN	1418	Ex	

	SaleCondition	SalePrice
523	Partial	184750




691	Normal	755000
1182	Abnorml	745000
1298	Partial	160000

- There are two houses with exceptionally high selling prices (over 700,000) and two with very low selling prices (under 200,000), all having a GrLivArea of more than 4,000.
- Interestingly, the houses sold at very low prices are actually larger (around 4,600 and 5,600 square feet). Upon closer inspection, the low-priced houses were sold under a partial sale condition, suggesting they were likely unfinished.
- In contrast, the high-priced house has a smaller area than the other two, possibly due to other factors such as house quality or location influencing the price.

## 2.4 d. Boxplot of OverallQual vs Target Variable (Sale Price)

```
[41]: data = pd.concat([df_train['SalePrice'], df_train['OverallQual']], axis=1)
f, ax = plt.subplots(figsize=(8, 6))
fig = sns.boxplot(x='OverallQual', y="SalePrice", data=data)
plt.show()
```



- As the OverallQual (represents the overall quality of the house) increases, the SalePrice also tends to increase.

- There are several outliers, particularly for high-quality houses (OverallQual of 9 and 10). For OverallQual = 10, there is an extreme outlier with a SalePrice exceeding 700,000 or under 200,000.

[43]: *# Check the anomaly or outliers:*

```
df_train[((df_train['OverallQual'] == 10) & (df_train['SalePrice'] < 200000)) |
          ((df_train['OverallQual'] == 10) & (df_train['SalePrice'] > 700000))]
```

[43]:

	LotArea	YearBuilt	OverallQual	OverallCond	RoofMatl	ExterQual	\
523	40094	2007	10	5	CompShg	Ex	
691	21535	1994	10	6	WdShngl	Ex	
1182	15623	1996	10	5	CompShg	Gd	
1298	63887	2008	10	5	ClyTile	Ex	

	BedroomAbvGr	GrLivArea	MiscFeature	GarageArea	KitchenQual	\
523	3	4676	NaN	884	Ex	
691	4	4316	NaN	832	Ex	
1182	4	4476	NaN	813	Ex	
1298	3	5642	NaN	1418	Ex	

	SaleCondition	SalePrice
523	Partial	184750
691	Normal	755000
1182	Abnorml	745000
1298	Partial	160000

- Houses with high OverallQual but low Sale Price (< 200,000) is houses with a partial sale condition, where the house is likely to be sold before completion and still under construction (unfinished work)
- Two other houses have high OverallQual and very high prices (>700,000): One of the houses was sold in normal conditions with an overall condition rating of 6 and features a high-quality luxury roof made of wood, which is aesthetic and eco-friendly, making it expensive. Meanwhile, the other house was sold under an abnormal sale condition, suggesting that it could be a luxury property or possess unique high-end features.

## 2.5 e. Boxplot of YearBuilt vs Target Variable (Sale Price)

[48]: `print(df_train['YearBuilt'].min(), df_train['YearBuilt'].max())`


1872 2010

The oldest house was built in 1872, while the newest one was built in 2010.

[51]: *# Check whether there is a correlation between the year of manufacture and the ↵  
↵ sale price:*

```
data = pd.concat([df_train['SalePrice'], df_train['YearBuilt']], axis=1)
```

```
f, ax = plt.subplots(figsize=(16, 8))
fig = sns.boxplot(x='YearBuilt', y='SalePrice', data=data)
plt.show()
```




There is no correlation between the year a house was built (Year Built) and Sale Price.

## 2.6 f. Correlation among Variables

```
[54]: # Correlation map among variables:

plt.figure(figsize=(10,10))
ax = sns.heatmap(df_train.select_dtypes(exclude = 'object').corr(), cmap =_
    ↪ 'coolwarm', annot=True, linewidth=2)
plt.show()
```

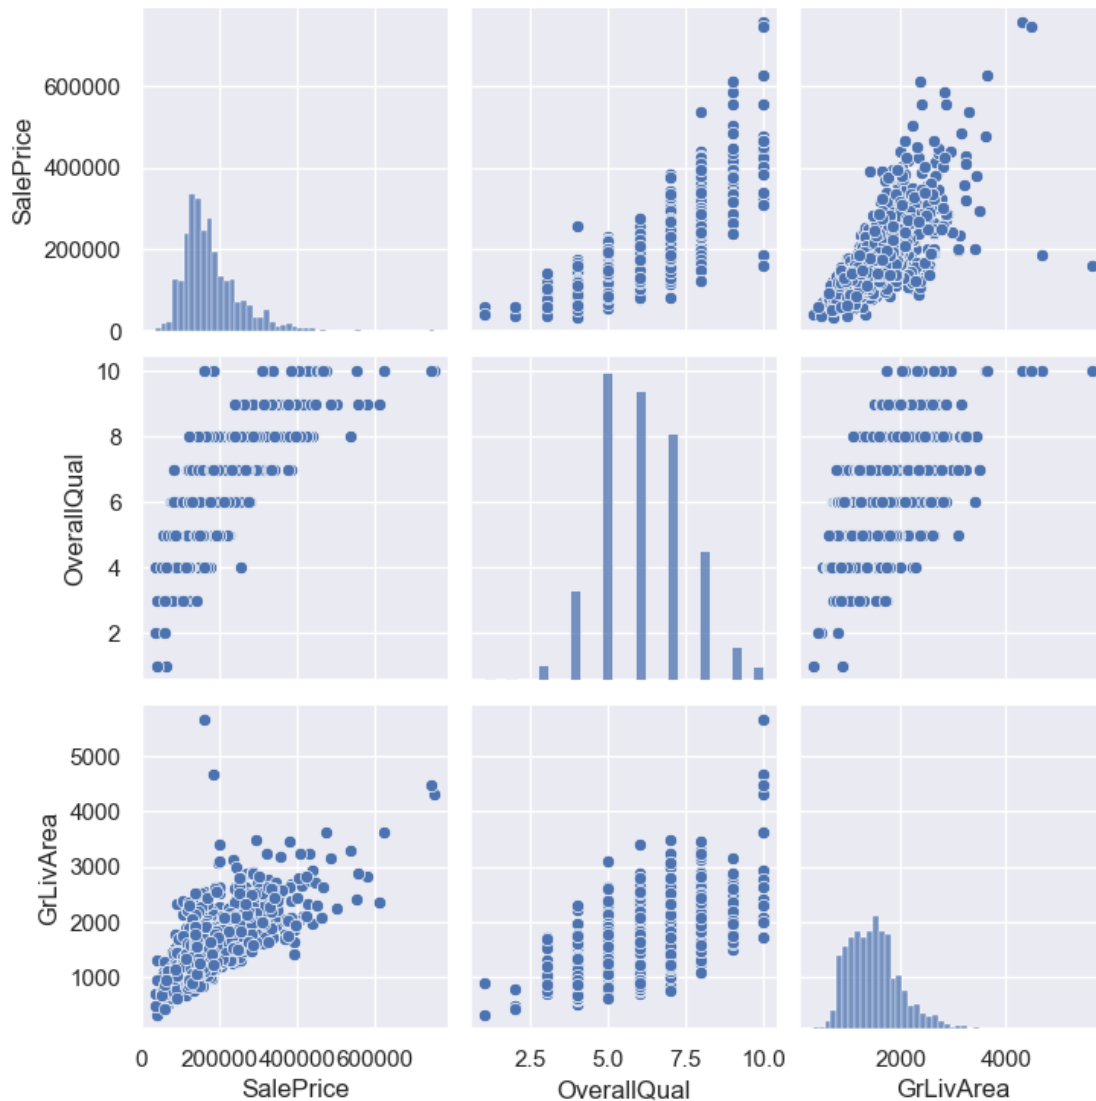


Insight: There is a positive correlation between Overall Quality (0.79) or GrLiving Area (0.71) and Sale Price.

```
[56]: # Pairplot between three variables:

sns.set()
cols = ['SalePrice', 'OverallQual', 'GrLivArea']
sns.pairplot(df_train[cols], size = 2.5)
```


```
plt.show()
```



- Strong positive correlation between Sale Price and Overall Quality and GrLiving Area.
- Moderate correlation between Overall Quality and GrLiving Area: higher quality houses tend to be larger, but size is not the only factor.
- Sale Price and GrLiving Area are right-skewed distributions.


## 2.7 g. Distribution of Target Variable (SalePrice column)

```
[59]: sns.distplot(df_train['SalePrice'], fit = norm)
plt.show()
```



*The sale prices is not normally distributed. So, preprocessing (transformation) is needed before further analysis.*

```
[62]: # Re-check the SalePrice using QQ-plot (Quantile-Quantile plot) whether the
      ↪ data follows a normal distribution:
fig = plt.figure()
res = stats.probplot(df_train['SalePrice'], plot=plt)
plt.show()
```



- The SalePrice is NOT normally distributed and is right-skewed with high-priced outliers.
- Log transformation can be applied to reduce this skewness.

### 3. Data Preprocessing

#### 3.1 a. Log Transformation for Target Variable (Sale Price)

```
[67]: df_train_tf = df_train.copy()      #copy the original dataframe

[68]: # Before applying the transformation to the entire dataset, we first test it on
      ↪ two data points, one of which is an outlier:

      # Before transformation:
      print(755000 - 180000)

      print("")

      # After transformation:
      print(np.log1p(755000))
```

```
print(np.log1p(180000))

print("")
print(np.log1p(755000) - np.log1p(180000))
```

575000

13.534474352733596


12.100717685412471

1.4337566673211253

- The difference between the two saleprices before the transformation is 575,000 (LARGE).
- The difference between the two saleprices after the transformation is 1.44 (SMALL).

```
[70]: # Apply Log Transformation to SalePrice column:
df_train_tf['SalePrice'] = np.log1p(df_train['SalePrice'])
```

```
[71]: # Re-check the distribution:
sns.distplot(df_train_tf['SalePrice'], fit = norm)
plt.show()
```





This confirms that the transformation effectively made SalePrice more normally distributed. Statistical models and machine learning algorithms typically perform better when the data follows a normal distribution.

### 3.2 b. Merge Train and Test Data

```
[74]: ntrain = df_train_tf.shape[0]
ntrain = df_train_tf.shape[0]
ntest = df_test.shape[0]
y_train = df_train_tf.SalePrice.values
all_data = pd.concat((df_train_tf, df_test)).reset_index(drop=True)
all_data.drop(['SalePrice'], axis=1, inplace=True)
print("all_data size is : {}".format(all_data.shape))
```

all\_data size is : (2919, 12)

```
[75]: all_data
```

```
[75]:
```

	LotArea	YearBuilt	OverallQual	OverallCond	RoofMatl	ExterQual	\
0	8450	2003	7	5	CompShg	Gd	
1	9600	1976	6	8	CompShg	TA	
2	11250	2001	7	5	CompShg	Gd	
3	9550	1915	7	5	CompShg	TA	
4	14260	2000	8	5	CompShg	Gd	
...	...	...	...	...	...	...	
2914	1936	1970	4	7	CompShg	TA	
2915	1894	1970	4	5	CompShg	TA	
2916	20000	1960	5	7	CompShg	TA	
2917	10441	1992	5	5	CompShg	TA	
2918	9627	1993	7	5	CompShg	TA	

	BedroomAbvGr	GrLivArea	MiscFeature	GarageArea	KitchenQual	\
0	3	1710	NaN	548.0	Gd	
1	3	1262	NaN	460.0	TA	
2	3	1786	NaN	608.0	Gd	
3	3	1717	NaN	642.0	Gd	
4	4	2198	NaN	836.0	Gd	
...	...	...	...	...	...	
2914	3	1092	NaN	0.0	TA	
2915	3	1092	NaN	286.0	TA	
2916	4	1224	NaN	576.0	TA	
2917	3	970	Shed	0.0	TA	
2918	3	2000	NaN	650.0	TA	

	SaleCondition
0	Normal
1	Normal
2	Normal

```

3      Abnorml
4      Normal
...
2914   Normal
2915   Abnorml
2916   Abnorml
2917   Normal
2918   Normal

```

[2919 rows x 12 columns]

[76]: `all_data.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2919 entries, 0 to 2918
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   LotArea         2919 non-null  int64
1   YearBuilt       2919 non-null  int64
2   OverallQual     2919 non-null  int64
3   OverallCond     2919 non-null  int64
4   RoofMatl       2919 non-null  object
5   ExterQual       2919 non-null  object
6   BedroomAbvGr   2919 non-null  int64
7   GrLivArea       2919 non-null  int64
8   MiscFeature     105 non-null   object
9   GarageArea      2918 non-null  float64
10  KitchenQual     2918 non-null  object
11  SaleCondition   2919 non-null  object
dtypes: float64(1), int64(6), object(5)
memory usage: 273.8+ KB

```

There are three columns that have missing values: MiscFeature, GarageArea, and KitchenQual.

### 3.3 c. Handling Missing Value Analysis and Imputation

[79]: *# Check the proportion of missing value in the dataset:*

```

total = all_data.isnull().sum().sort_values(ascending=False)
percent = (all_data.isnull().sum()/all_data.isnull().count() * 100).
    ↪sort_values(ascending=False)
missing_data = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
missing_data

```

[79]:

	Total	Percent
MiscFeature	2814	96.402878
GarageArea	1	0.034258

KitchenQual	1	0.034258
LotArea	0	0.000000
YearBuilt	0	0.000000
OverallQual	0	0.000000
OverallCond	0	0.000000
RoofMatl	0	0.000000
ExterQual	0	0.000000
BedroomAbvGr	0	0.000000
GrLivArea	0	0.000000
SaleCondition	0	0.000000

- Miscfeature has 2814 missing values.
- Garage Area and Kitchen Quality have 1 missing value.

### 3.3.1 MiscFeature column

```
[83]: # Count each unique value in MiscFeature column:
all_data['MiscFeature'].value_counts()
```

```
[83]: MiscFeature
Shed      95
Gar2       5
Othr       4
TenC       1
Name: count, dtype: int64
```

```
[85]: # The rows with null values are filled with "NA":
all_data['MiscFeature'] = all_data['MiscFeature'].fillna('NA')
```

### 3.3.2 GarageArea column

```
[89]: all_data['GarageArea'].value_counts()
```

```
[89]: GarageArea
0.0      157
576.0     97
440.0     96
240.0     69
484.0     68
...
872.0      1
923.0      1
192.0      1
1025.0     1
272.0      1
Name: count, Length: 603, dtype: int64
```

```
[92]: # The rows with null values are filled with 0 (No Garage):
all_data['GarageArea'] = all_data['GarageArea'].fillna(0)
```

### 3.3.3 KitchenQual column

```
[95]: all_data['KitchenQual'].value_counts()
```

```
[95]: KitchenQual
TA      1492
Gd      1151
Ex       205
Fa        70
Name: count, dtype: int64
```

```
[100]: # The rows with null values are filled with mode (the most frequently occurring
↪value):
all_data['KitchenQual'] = all_data['KitchenQual'].
↪fillna(all_data['KitchenQual'].mode()[0])
```

```
[103]: # Check the remaining missing values (if any):
all_data_na = (all_data.isnull().sum() / len(all_data)) * 100
all_data_na = all_data_na.drop(all_data_na[all_data_na == 0].index).
↪sort_values(ascending=False)
missing_data = pd.DataFrame({'Missing Ratio' :all_data_na})
missing_data.head()
```

```
[103]: Empty DataFrame
Columns: [Missing Ratio]
Index: []
```

```
[106]: all_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2919 entries, 0 to 2918
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   LotArea          2919 non-null  int64
1   YearBuilt        2919 non-null  int64
2   OverallQual      2919 non-null  int64
3   OverallCond      2919 non-null  int64
4   RoofMatl         2919 non-null  object
5   ExterQual        2919 non-null  object
6   BedroomAbvGr     2919 non-null  int64
7   GrLivArea        2919 non-null  int64
8   MiscFeature      2919 non-null  object
9   GarageArea       2919 non-null  float64
```

```

10 KitchenQual      2919 non-null    object
11 SaleCondition    2919 non-null    object
dtypes: float64(1), int64(6), object(5)
memory usage: 273.8+ KB

```

There is no missing value. All data is complete.

### 3.4 d. Change the Categorical Data Type

```

[113]: # Change the data type to string:
all_data['OverallQual'] = all_data['OverallQual'].astype(str)
all_data['OverallCond'] = all_data['OverallCond'].astype(str)

```

```

[115]: all_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2919 entries, 0 to 2918
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   LotArea         2919 non-null  int64
1   YearBuilt       2919 non-null  int64
2   OverallQual     2919 non-null  object
3   OverallCond     2919 non-null  object
4   RoofMatl       2919 non-null  object
5   ExterQual       2919 non-null  object
6   BedroomAbvGr   2919 non-null  int64
7   GrLivArea      2919 non-null  int64
8   MiscFeature     2919 non-null  object
9   GarageArea     2919 non-null  float64
10  KitchenQual     2919 non-null  object
11  SaleCondition   2919 non-null  object
dtypes: float64(1), int64(4), object(7)
memory usage: 273.8+ KB

```

- The data type is changed from integer to object/string because the data is in the form of ordinal categorical, not numerical.
- After being converted to object/string, label encoding is applied before processing with machine learning

### 3.5 e. Label Encoding

```

[119]: tmp_data = all_data.copy()

```

```

[121]: from sklearn.preprocessing import LabelEncoder
cols = ('OverallQual', 'OverallCond', 'ExterQual', 'KitchenQual') # ordinal_
      ↪ categorical data
encoders = {}

```

```

for c in cols:
    lbl = LabelEncoder() # create encoder objects
    lbl.fit(list(tmp_data[c].values)) # encoder objects learn their data
    tmp_data[c] = lbl.transform(list(tmp_data[c].values)) # encoder object
    ↪change their data
    encoders[c] = lbl # save each encoder for deployment

# shape
print('Shape tmp_data: {}'.format(tmp_data.shape))

```

Shape tmp\_data: (2919, 12)

[122]: tmp\_data

```

[122]:
   LotArea  YearBuilt  OverallQual  OverallCond  RoofMatl  ExterQual  \
0      8450      2003           7           4  CompShg           2
1      9600      1976           6           7  CompShg           3
2     11250      2001           7           4  CompShg           2
3      9550      1915           7           4  CompShg           3
4     14260      2000           8           4  CompShg           2
...      ...      ...      ...      ...      ...      ...
2914     1936      1970           4           6  CompShg           3
2915     1894      1970           4           4  CompShg           3
2916    20000      1960           5           6  CompShg           3
2917    10441      1992           5           4  CompShg           3
2918     9627      1993           7           4  CompShg           3

   BedroomAbvGr  GrLivArea  MiscFeature  GarageArea  KitchenQual  \
0              3      1710           NA      548.0           2
1              3      1262           NA      460.0           3
2              3      1786           NA      608.0           2
3              3      1717           NA      642.0           2
4              4      2198           NA      836.0           2
...      ...      ...      ...      ...      ...
2914              3      1092           NA           0.0           3
2915              3      1092           NA      286.0           3
2916              4      1224           NA      576.0           3
2917              3       970        Shed           0.0           3
2918              3      2000           NA      650.0           3

   SaleCondition
0      Normal
1      Normal
2      Normal
3    Abnorml
4      Normal

```

```

...
2914      Normal
2915      Abnorml
2916      Abnorml
2917      Normal
2918      Normal

[2919 rows x 12 columns]

```

### 3.6 f. Distribution analysis for feature columns (to see skewness)

```

[126]: from scipy.stats import norm, skew

# for numerical columns:
numeric_feats = tmp_data.dtypes[tmp_data.dtypes != "object"].index

# Check the skew of all numerical features
skewed_feats = tmp_data[numeric_feats].apply(lambda x: skew(x.dropna())).
↳sort_values(ascending=False)
print("\nSkewness in numerical features: \n")
skewness = pd.DataFrame({'Skew' :skewed_feats})
skewness

```

Skewness in numerical features:

```

[126]:


```

	Skew
LotArea	12.822431
GrLivArea	1.269358
OverallCond	0.570312
BedroomAbvGr	0.326324
GarageArea	0.239257
OverallQual	-0.326653
YearBuilt	-0.599806
KitchenQual	-1.448023
ExterQual	-1.801409

- Lot Area is highly right-skewed, so it needs to be transformed to prevent the model from underfitting.
- Meanwhile, GrLivArea, KitchenQual, and ExterQual which have moderate skewness, also need to be transformed to prevent the model from overfitting.

### 3.6.1 Lot Area

```
[130]: # lot Area column distribution:  
sns.distplot(tmp_data['LotArea'])  
plt.show()
```



```
[131]: tmp_data['LotArea'].describe()
```

```
[131]: count      2919.000000  
mean       10168.114080  
std         7886.996359  
min          1300.000000  
25%         7478.000000  
50%         9453.000000  
75%        11570.000000  
max        215245.000000  
Name: LotArea, dtype: float64
```

The smallest lot area is 1,300, with a maximum of up to 215,245, and an average of 10,168.



### 3.7 g. Box Cox Transformation for feature columns with skewness >0.75 atau <-0.75

```
[136]: skewness = skewness[abs(skewness['Skew']) > 0.75]
print("There are {} skewed numerical features to Box Cox transform".
      ↪format(skewness.shape[0]))
```

There are 4 skewed numerical features to Box Cox transform

```
[138]: # Box Cox transformation:
from scipy.special import boxcox1p

skewed_features = skewness.index
lam = 0.20
for feat in skewed_features:
    tmp_data[feat] = boxcox1p(tmp_data[feat], lam)
```

```
[140]: # features to be transformed:
skewed_features
```

```
[140]: Index(['LotArea', 'GrLivArea', 'KitchenQual', 'ExterQual'], dtype='object')
```

```
[142]: # results:
tmp_data
```

```
[142]:
```

	LotArea	YearBuilt	OverallQual	OverallCond	RoofMatl	ExterQual	\
0	25.503637	2003	7	4	CompShg	1.228655	
1	26.291998	1976	6	7	CompShg	1.597540	
2	27.300424	2001	7	4	CompShg	1.228655	
3	26.259338	1915	7	4	CompShg	1.597540	
4	28.868815	2000	8	4	CompShg	1.228655	
...	...	...	...	...	...	...	
2914	17.719351	1970	4	6	CompShg	1.597540	
2915	17.619961	1970	4	4	CompShg	1.597540	
2916	31.239346	1960	5	6	CompShg	1.597540	
2917	26.821947	1992	5	4	CompShg	1.597540	
2918	26.309578	1993	7	4	CompShg	1.597540	

	BedroomAbvGr	GrLivArea	MiscFeature	GarageArea	KitchenQual	\
0	3	17.162564	NA	548.0	1.228655	
1	3	15.856944	NA	460.0	1.597540	
2	3	17.356042	NA	608.0	1.228655	
3	3	17.180669	NA	642.0	1.228655	
4	4	18.303173	NA	836.0	1.228655	
...	...	...	...	...	...	
2914	3	15.262547	NA	0.0	1.597540	
2915	3	15.262547	NA	286.0	1.597540	
2916	4	15.729901	NA	576.0	1.597540	

2917	3	14.788544	Shed	0.0	1.597540
2918	3	17.867539	NA	650.0	1.597540

	SaleCondition
0	Normal
1	Normal
2	Normal
3	Abnorml
4	Normal
...	...
2914	Normal
2915	Abnorml
2916	Abnorml
2917	Normal
2918	Normal

[2919 rows x 12 columns]

```
[144]: # Check the skewness of all numerical features:
skewed_feats = tmp_data[numeric_feats].apply(lambda x: skew(x.dropna())).
    ↪sort_values(ascending=False)
print("\nSkew in numerical features: \n")
skewness = pd.DataFrame({'Skew' :skewed_feats})
skewness
```

Skew in numerical features:

```
[144]:
```

	Skew
OverallCond	0.570312
LotArea	0.496692
BedroomAbvGr	0.326324
GarageArea	0.239257
GrLivArea	0.230000
OverallQual	-0.326653
YearBuilt	-0.599806
KitchenQual	-2.156088
ExterQual	-2.778802

### 3.8 h. One-Hot Encoding for non-ordinal categorical columns

```
[147]: tmp_data
```

```
[147]:
```

	LotArea	YearBuilt	OverallQual	OverallCond	RoofMatl	ExterQual	\
0	25.503637	2003	7	4	CompShg	1.228655	
1	26.291998	1976	6	7	CompShg	1.597540	

2	27.300424	2001	7	4	CompShg	1.228655
3	26.259338	1915	7	4	CompShg	1.597540
4	28.868815	2000	8	4	CompShg	1.228655
...	...	...	...	...	...	...
2914	17.719351	1970	4	6	CompShg	1.597540
2915	17.619961	1970	4	4	CompShg	1.597540
2916	31.239346	1960	5	6	CompShg	1.597540
2917	26.821947	1992	5	4	CompShg	1.597540
2918	26.309578	1993	7	4	CompShg	1.597540

	BedroomAbvGr	GrLivArea	MiscFeature	GarageArea	KitchenQual	\
0	3	17.162564	NA	548.0	1.228655	
1	3	15.856944	NA	460.0	1.597540	
2	3	17.356042	NA	608.0	1.228655	
3	3	17.180669	NA	642.0	1.228655	
4	4	18.303173	NA	836.0	1.228655	
...	...	...	...	...	...	
2914	3	15.262547	NA	0.0	1.597540	
2915	3	15.262547	NA	286.0	1.597540	
2916	4	15.729901	NA	576.0	1.597540	
2917	3	14.788544	Shed	0.0	1.597540	
2918	3	17.867539	NA	650.0	1.597540	

	SaleCondition
0	Normal
1	Normal
2	Normal
3	Abnorml
4	Normal
...	...
2914	Normal
2915	Abnorml
2916	Abnorml
2917	Normal
2918	Normal

[2919 rows x 12 columns]

```
[149]: tmp_data = pd.get_dummies(tmp_data, drop_first=True) # one hot encoding for
      ↪ other categorical data
      print(tmp_data.shape)
```

(2919, 25)

```
[151]: # Check the results:
      tmp_data
```

[151]:

	LotArea	YearBuilt	OverallQual	OverallCond	ExterQual	BedroomAbvGr	\
0	25.503637	2003	7	4	1.228655	3	
1	26.291998	1976	6	7	1.597540	3	
2	27.300424	2001	7	4	1.228655	3	
3	26.259338	1915	7	4	1.597540	3	
4	28.868815	2000	8	4	1.228655	4	
...	...	...	...	...	...	...	
2914	17.719351	1970	4	6	1.597540	3	
2915	17.619961	1970	4	4	1.597540	3	
2916	31.239346	1960	5	6	1.597540	4	
2917	26.821947	1992	5	4	1.597540	3	
2918	26.309578	1993	7	4	1.597540	3	

	GrLivArea	GarageArea	KitchenQual	RoofMatl_CompShg	...	\
0	17.162564	548.0	1.228655	True	...	
1	15.856944	460.0	1.597540	True	...	
2	17.356042	608.0	1.228655	True	...	
3	17.180669	642.0	1.228655	True	...	
4	18.303173	836.0	1.228655	True	...	
...	...	...	...	...	...	
2914	15.262547	0.0	1.597540	True	...	
2915	15.262547	286.0	1.597540	True	...	
2916	15.729901	576.0	1.597540	True	...	
2917	14.788544	0.0	1.597540	True	...	
2918	17.867539	650.0	1.597540	True	...	

	RoofMatl_WdShngl	MiscFeature_NA	MiscFeature_Othr	MiscFeature_Shed	\
0	False	True	False	False	
1	False	True	False	False	
2	False	True	False	False	
3	False	True	False	False	
4	False	True	False	False	
...	...	...	...	...	
2914	False	True	False	False	
2915	False	True	False	False	
2916	False	True	False	False	
2917	False	False	False	True	
2918	False	True	False	False	

	MiscFeature_TenC	SaleCondition_AdjLand	SaleCondition_Alloca	\
0	False	False	False	
1	False	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	False	
...	...	...	...	
2914	False	False	False	

2915	False	False	False
2916	False	False	False
2917	False	False	False
2918	False	False	False

	SaleCondition_Family	SaleCondition_Normal	SaleCondition_Partial
0	False	True	False
1	False	True	False
2	False	True	False
3	False	False	False
4	False	True	False
...	...	...	...
2914	False	True	False
2915	False	False	False
2916	False	False	False
2917	False	True	False
2918	False	True	False

[2919 rows x 25 columns]

There is an additional columns from 12 to 25 (from RoofMatl, MiscFeature, and SaleCondition columns)

```
[154]: # save the encoding result columns for later use:
one_hot_columns = tmp_data.columns
```

### 3.9 i. Scaling

```
[157]: # Check the value range:
tmp_data
```

```
[157]:
```

	LotArea	YearBuilt	OverallQual	OverallCond	ExterQual	BedroomAbvGr	\
0	25.503637	2003	7	4	1.228655	3	
1	26.291998	1976	6	7	1.597540	3	
2	27.300424	2001	7	4	1.228655	3	
3	26.259338	1915	7	4	1.597540	3	
4	28.868815	2000	8	4	1.228655	4	
...	...	...	...	...	...	...	
2914	17.719351	1970	4	6	1.597540	3	
2915	17.619961	1970	4	4	1.597540	3	
2916	31.239346	1960	5	6	1.597540	4	
2917	26.821947	1992	5	4	1.597540	3	
2918	26.309578	1993	7	4	1.597540	3	

	GrLivArea	GarageArea	KitchenQual	RoofMatl_CompShg	...	\
0	17.162564	548.0	1.228655	True	...	
1	15.856944	460.0	1.597540	True	...	

2	17.356042	608.0	1.228655	True	...
3	17.180669	642.0	1.228655	True	...
4	18.303173	836.0	1.228655	True	...
...	...	...	...	...	...
2914	15.262547	0.0	1.597540	True	...
2915	15.262547	286.0	1.597540	True	...
2916	15.729901	576.0	1.597540	True	...
2917	14.788544	0.0	1.597540	True	...
2918	17.867539	650.0	1.597540	True	...

	RoofMatl_WdShngl	MiscFeature_NA	MiscFeature_Othr	MiscFeature_Shed	\
0	False	True	False	False	
1	False	True	False	False	
2	False	True	False	False	
3	False	True	False	False	
4	False	True	False	False	
...	...	...	...	...	
2914	False	True	False	False	
2915	False	True	False	False	
2916	False	True	False	False	
2917	False	False	False	True	
2918	False	True	False	False	

	MiscFeature_TenC	SaleCondition_AdjLand	SaleCondition_Alloca	\
0	False	False	False	
1	False	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	False	
...	...	...	...	
2914	False	False	False	
2915	False	False	False	
2916	False	False	False	
2917	False	False	False	
2918	False	False	False	

	SaleCondition_Family	SaleCondition_Normal	SaleCondition_Partial
0	False	True	False
1	False	True	False
2	False	True	False
3	False	False	False
4	False	True	False
...	...	...	...
2914	False	True	False
2915	False	False	False
2916	False	False	False
2917	False	True	False

2918                      False                      True                      False

[2919 rows x 25 columns]

```
[159]: # Use a robust scaler to reduce the impact of outliers in each column (far
↳range of values):
from sklearn.preprocessing import RobustScaler
scaler = RobustScaler()
scaler.fit(tmp_data)
tmp_data = scaler.transform(tmp_data)
```

```
[161]: tmp_data
```

```
[161]: array([[ -0.25487439,  0.63157895,  0.5          , ...,  0.          ,
          0.          ,  0.          ],
        [ 0.03551351,  0.06315789,  0.          , ...,  0.          ,
          0.          ,  0.          ],
        [ 0.40696091,  0.58947368,  0.5          , ...,  0.          ,
          0.          ,  0.          ],
        ...,
        [ 1.85783766, -0.27368421, -0.5          , ...,  0.          ,
        -1.          ,  0.          ],
        [ 0.2307166 ,  0.4          , -0.5          , ...,  0.          ,
          0.          ,  0.          ],
        [ 0.04198905,  0.42105263,  0.5          , ...,  0.          ,
          0.          ,  0.          ]])
```

## 4 4. Split Data Training dan Testing

```
[164]: X_train = tmp_data[:ntrain]
X_test = tmp_data[ntrain:]

X_train.shape, X_test.shape, y_train.shape
```

```
[164]: ((1460, 25), (1459, 25), (1460,))
```

## 5 6. Modelling

```
[167]: from sklearn.model_selection import KFold, cross_val_score
from sklearn.metrics import make_scorer, r2_score

# Cross validation:
def test_model(model, X_train=X_train, y_train=y_train):    # create a function
    cv = KFold(n_splits = 4, shuffle=True, random_state = 45)    # 4x cross
↳validation, shuffle
```

```

r2 = make_scorer(r2_score)

r2_val_score = cross_val_score(model, X_train, y_train, cv=cv, scoring = 'r2') # calculate R2
score = [r2_val_score.mean()]
return score

```

## 5.1 a. Linear Regression

```

[170]: import sklearn.linear_model as linear_model
LR = linear_model.LinearRegression()
test_model(LR)

```

[170]: [0.8240757196006232]

## 5.2 b. Lasso Regression

```

[173]: lasso = linear_model.Lasso(alpha=1e-4)
test_model(lasso)

```

[173]: [0.8335074727025339]

## 5.3 c. Random Forest

```

[176]: random_forest = linear_model.Ridge(alpha=1)
test_model(random_forest)

```

[176]: [0.8363080677683358]

Lasso and Random Forest Regression show better regression compared to others

## 5.4 d. Support Vector Regression

```

[180]: from sklearn.svm import SVR
svr_reg = SVR(kernel= 'rbf')
test_model(svr_reg)

```

[180]: [0.8245841296419294]

## 5.5 e. XGBoost

```

[182]: !pip install --upgrade xgboost

```

Requirement already satisfied: xgboost in c:\users\62817\anaconda3\lib\site-packages (3.0.0)

Requirement already satisfied: numpy in c:\users\62817\anaconda3\lib\site-



packages (from xgboost) (1.26.4)  
Requirement already satisfied: scipy in c:\users\62817\anaconda3\lib\site-  
packages (from xgboost) (1.13.1)

```
[184]: import xgboost
```

```
xgb_reg = xgboost.XGBRegressor()  
test_model(xgb_reg)
```

```
[184]: [0.815044432921986]
```

## 5.6 f. Optimize the model: Lasso Regression vs Random Forest

### Lasso Regression

```
[187]: from sklearn.model_selection import GridSearchCV
```

```
param_grid = {'alpha': [1e-4, 1e-3, 1e-2, 0.1, 1, 10]}  
grid_search = GridSearchCV(linear_model.Lasso(), param_grid, cv=5)  
grid_search.fit(X_train, y_train)  
  
print(f"Best Alpha: {grid_search.best_params_['alpha']}")
```

Best Alpha: 0.001

```
[188]: from sklearn.linear_model import Lasso
```

```
lasso_best = Lasso(alpha=0.001)  
lasso_best.fit(X_train, y_train)
```

```
[188]: Lasso(alpha=0.001)
```

```
[189]: from sklearn.model_selection import cross_val_score
```

```
cv_scores = cross_val_score(lasso_best, X_train, y_train, cv=5, scoring='r2')  
  
print(f"Cross-Validation R2 Scores: {cv_scores}")  
print(f"Mean R2 Score: {cv_scores.mean():.4f}")  
print(f"Standard Deviation: {cv_scores.std():.4f}")
```

Cross-Validation R<sup>2</sup> Scores: [0.87173941 0.85853769 0.83866367 0.86223495  
0.78697472]

Mean R<sup>2</sup> Score: 0.8436

Standard Deviation: 0.0303

```
[190]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt
```

```

from sklearn.linear_model import Lasso

# Train Lasso model:
lasso = Lasso(alpha=0.001)
lasso.fit(X_train, y_train)

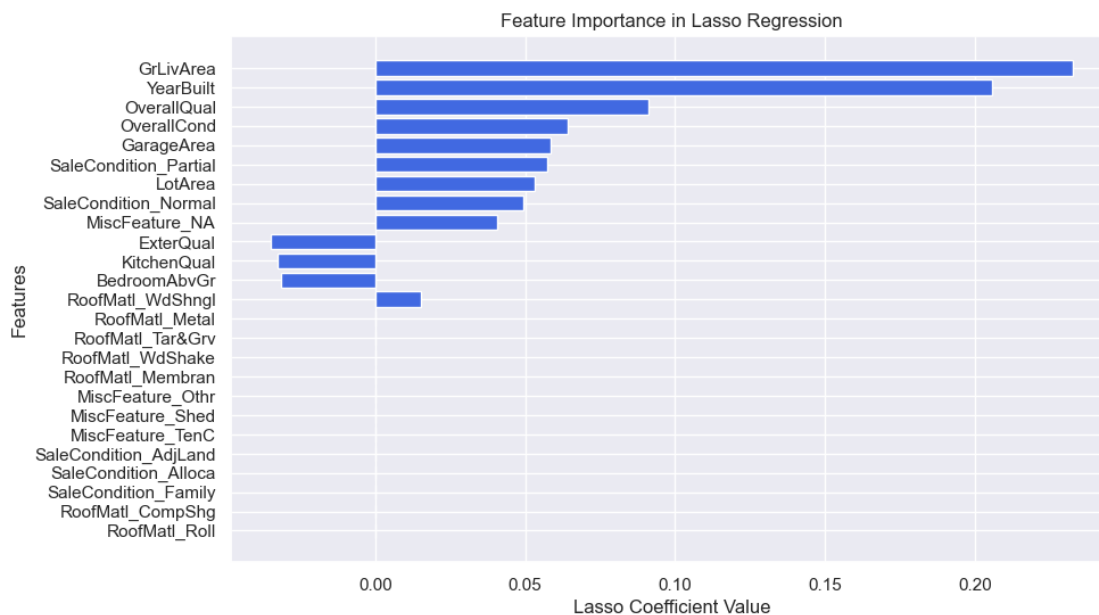
# create feature coefficients:
lasso_coefs = lasso.coef_

# Create DataFrame for visualization:
# Replace X_train.columns with one_hot_columns
lasso_importance_df = pd.DataFrame({'Feature': one_hot_columns, 'Importance':
    ↳lasso_coefs})
lasso_importance_df = lasso_importance_df.sort_values(by='Importance',
    ↳ascending=False, key=abs) # Sort by absolute value

# Visualization:
plt.figure(figsize=(10,6))
plt.barh(lasso_importance_df['Feature'], lasso_importance_df['Importance'],
    ↳color='royalblue')
plt.xlabel('Lasso Coefficient Value')
plt.ylabel('Features')
plt.title('Feature Importance in Lasso Regression')
plt.gca().invert_yaxis()
plt.show()

# Show the most influential features:
print(lasso_importance_df.head(10)) # Top 10 the important features

```



	Feature	Importance
6	GrLivArea	0.232665
1	YearBuilt	0.205813
2	OverallQual	0.091143
3	OverallCond	0.064278
7	GarageArea	0.058531
24	SaleCondition_Partial	0.057317
0	LotArea	0.053069
23	SaleCondition_Normal	0.049281
16	MiscFeature_NA	0.040732
4	ExterQual	-0.034696

### Random Forest

```
[192]: from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestRegressor

param_grid = {
    'n_estimators': [100, 300, 500],
    'max_depth': [10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt']
}

rf = RandomForestRegressor()
grid_search = GridSearchCV(rf, param_grid, cv=5, scoring='r2', n_jobs=-1)
grid_search.fit(X_train, y_train)

# Best parameters
print(grid_search.best_params_)
```

```
{'max_depth': 30, 'max_features': 'sqrt', 'min_samples_leaf': 1,
'min_samples_split': 2, 'n_estimators': 300}
```

```
[193]: from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestRegressor

# Use the best parameters from GridSearchCV
best_rf = RandomForestRegressor(
    max_depth=30,
    max_features='sqrt',
    min_samples_leaf=1,
    min_samples_split=2,
    n_estimators=500,
    random_state=42)
```

```
)

# Perform cross-validation
scores = cross_val_score(best_rf, X_train, y_train, cv=5, scoring='r2')

print("Mean Cross-Validation R2:", scores.mean())
print("Standard Deviation:", scores.std())
```

Mean Cross-Validation R<sup>2</sup>: 0.8500130283032352  
Standard Deviation: 0.011010251796251471

```
[194]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestRegressor

# Let X_train be the features and y_train be the target
rf = RandomForestRegressor(n_estimators=500, random_state=42)
rf.fit(X_train, y_train)

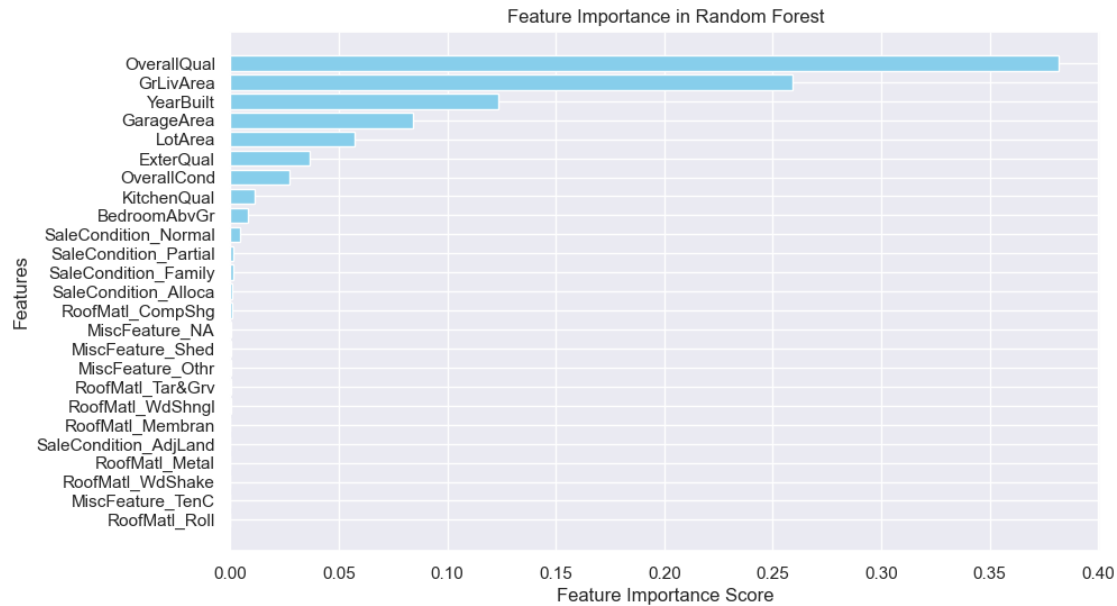
# Get the important features:
importances = rf.feature_importances_

# **Get feature names from the training data**
# Assuming one_hot_columns contains the feature names from before the scaling
↳step
feature_names = one_hot_columns

# Create DataFrame for visualization:
feature_importance_df = pd.DataFrame({'Feature': feature_names, 'Importance':
↳importances})
feature_importance_df = feature_importance_df.sort_values(by='Importance',
↳ascending=False)

# Show the results:
plt.figure(figsize=(10, 6))
plt.barh(feature_importance_df['Feature'], feature_importance_df['Importance'],
↳color='skyblue')
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.title('Feature Importance in Random Forest')
plt.gca().invert_yaxis()
plt.show()

print(feature_importance_df.head(10)) # Top 10 the important features
```



	Feature	Importance
2	OverallQual	0.381875
6	GrLivArea	0.259145
1	YearBuilt	0.123283
7	GarageArea	0.084152
0	LotArea	0.057279
4	ExterQual	0.036430
3	OverallCond	0.027518
8	KitchenQual	0.011399
5	BedroomAbvGr	0.008284
23	SaleCondition_Normal	0.004296

Based on the cross-validation results, both models perform well, but: - Random Forest has a slightly higher mean R2 score (0.8496 vs. 0.8436 for Lasso), indicating it might generalize slightly better. - Additionally, the standard deviation for Random Forest is lower (0.0142 vs. 0.0303), suggesting it is more stable across different validation splits.

Thus, since after tuning Random Forest has higher and more stable R<sup>2</sup>, it is the best choice to predict SalePrice.

## 6 7. Try to predict new data with the Lasso regression model (because the r2 score is the highest compared to other algorithms)

```
[198]: pd.DataFrame(tmp_data)      #created in dataframe form
```

```
[198]:
```

	0	1	2	3	4	5	6	7	8	9	\
0	-0.254874	0.631579	0.5	0.0	-1.0	0.0	0.395595	0.265625	-1.0	0.0	

```

1      0.035514  0.063158  0.0  3.0  0.0  0.0 -0.305745 -0.078125  0.0  0.0
2      0.406961  0.589474  0.5  0.0 -1.0  0.0  0.499525  0.500000 -1.0  0.0
3      0.023483 -1.221053  0.5  0.0  0.0  0.0  0.405320  0.632812 -1.0  0.0
4      0.984668  0.568421  1.0  0.0 -1.0  1.0  1.008295  1.390625 -1.0  0.0
...
2914 -3.122167 -0.063158 -1.0  2.0  0.0  0.0 -0.625037 -1.875000  0.0  0.0
2915 -3.158777 -0.063158 -1.0  0.0  0.0  0.0 -0.625037 -0.757812  0.0  0.0
2916  1.857838 -0.273684 -0.5  2.0  0.0  1.0 -0.373989  0.375000  0.0  0.0
2917  0.230717  0.400000 -0.5  0.0  0.0  0.0 -0.879657 -1.875000  0.0  0.0
2918  0.041989  0.421053  0.5  0.0  0.0  0.0  0.774286  0.664062  0.0  0.0

```

```

... 15 16 17 18 19 20 21 22 23 24
0 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
1 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
2 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
3 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 -1.0 0.0
4 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
...
2914 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
2915 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 -1.0 0.0
2916 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 -1.0 0.0
2917 ... 0.0 -1.0 0.0 1.0 0.0 0.0 0.0 0.0 0.0 0.0
2918 ... 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0

```

[2919 rows x 25 columns]

```

[199]: # random forest model:
from sklearn.ensemble import RandomForestRegressor

random_forest = RandomForestRegressor(n_estimators=500, random_state=42)

# Training the model with training data:
model = random_forest.fit(X_train, y_train)

```

```

[200]: tmp_data[0]

```

```

[200]: array([-0.25487439,  0.63157895,  0.5          ,  0.          , -1.          ,
              0.          ,  0.39559454,  0.265625   , -1.          ,  0.          ,
              0.          ,  0.          ,  0.          ,  0.          ,  0.          ,
              0.          ,  0.          ,  0.          ,  0.          ,  0.          ])

```

```

[201]: # Predict data to -0
model.predict(tmp_data[0].reshape(1,-1))

```

```

[201]: array([12.23591967])

```

```
[202]: # Predicted sale price:
np.expm1(12.25771747)
```

```
[202]: 210599.33513183476
```

```
[207]: # compare the predicted sale price with the actual house price:
y_train
```

```
[207]: array([12.24769912, 12.10901644, 12.31717117, ..., 12.49313327,
        11.86446927, 11.90159023])
```

```
[208]: # The actual sale price for first row houses in training data:
np.expm1(12.24769912)
```

```
[208]: 208500.00075632462
```

```
[209]: # try to input new data to predict how much the house will cost
new_data = {'LotArea': [8000],
            'YearBuilt': [2010],
            'OverallQual': ['8'],
            'OverallCond': ['7'],
            'GrLivArea': [2000],
            'MiscFeature': ['Shed'],
            'ExterQual': ['Gd'],
            'RoofMatl': ['Compshg'],
            'SaleCondition': ['Normal'],
            'GarageArea': [500],
            'KitchenQual': ['Gd'],}
```

```
[211]: tmp = pd.DataFrame(new_data)
tmp
```

```
[211]:   LotArea  YearBuilt  OverallQual  OverallCond  GrLivArea  MiscFeature \
0     8000      2010           8           7      2000      Shed

   ExterQual  RoofMatl  SaleCondition  GarageArea  KitchenQual
0         Gd  Compshg         Normal         500           Gd
```

```
[212]: new_data_tf = new_data.copy()
```

```
[213]: # label encoding for new data
new_data_tf['OverallQual'] = encoders['OverallQual'].
    ↪transform(tmp['OverallQual'])[0]
new_data_tf['OverallCond'] = encoders['OverallCond'].
    ↪transform(tmp['OverallCond'])[0]
new_data_tf['ExterQual'] = encoders['ExterQual'].transform(tmp['ExterQual'])[0]
```

```
new_data_tf['KitchenQual'] = encoders['KitchenQual'].
↳transform(tmp['KitchenQual'])[0]
```

```
[214]: # transformasi boxcox untuk kolom LotArea, GrLivArea, dan KitchenQual
new_data_tf['LotArea'] = boxcox1p(tmp['LotArea'], lam)[0]
new_data_tf['GrLivArea'] = boxcox1p(tmp['GrLivArea'], lam)[0]
new_data_tf['KitchenQual'] = boxcox1p(new_data_tf['KitchenQual'], lam)
new_data_tf['ExterQual'] = boxcox1p(new_data_tf['ExterQual'], lam)
```

```
[215]: tmp = pd.DataFrame(new_data_tf)
```

```
[216]: tmp = pd.get_dummies(tmp, columns = ['MiscFeature', 'RoofMatl',
↳'SaleCondition'])
tmp
```

```
[216]:      LotArea  YearBuilt  OverallQual  OverallCond  GrLivArea  ExterQual  \
0  25.171636      2010           8           6  17.867539  1.228655

      GarageArea  KitchenQual  MiscFeature_Shed  RoofMatl_Compshg  \
0           500      1.228655           True           True

      SaleCondition_Normal
0           True
```

```
[217]: # Adding missing columns to new data:
for kolom in one_hot_columns:
    if kolom not in tmp.columns:
        tmp[kolom] = 0
```

```
[218]: tmp = tmp.reindex(columns=one_hot_columns, fill_value=0)
```

```
[219]: tmp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1 entries, 0 to 0
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   LotArea                1 non-null     float64
1   YearBuilt              1 non-null     int64
2   OverallQual            1 non-null     int32
3   OverallCond            1 non-null     int32
4   ExterQual              1 non-null     float64
5   BedroomAbvGr          1 non-null     int64
6   GrLivArea              1 non-null     float64
7   GarageArea             1 non-null     int64
8   KitchenQual            1 non-null     float64
```



```

9   RoofMatl_CompShg      1 non-null    int64
10  RoofMatl_Membran      1 non-null    int64
11  RoofMatl_Metal        1 non-null    int64
12  RoofMatl_Roll         1 non-null    int64
13  RoofMatl_Tar&Grv      1 non-null    int64
14  RoofMatl_WdShake      1 non-null    int64
15  RoofMatl_WdShngl      1 non-null    int64
16  MiscFeature_NA        1 non-null    int64
17  MiscFeature_Othr      1 non-null    int64
18  MiscFeature_Shed      1 non-null    bool
19  MiscFeature_TenC      1 non-null    int64
20  SaleCondition_AdjLand  1 non-null    int64
21  SaleCondition_Alloca  1 non-null    int64
22  SaleCondition_Family  1 non-null    int64
23  SaleCondition_Normal  1 non-null    bool
24  SaleCondition_Partial 1 non-null    int64
dtypes: bool(2), float64(4), int32(2), int64(17)
memory usage: 310.0 bytes

```

```
[220]: # Scaling
tmp_scaled = scaler.transform(tmp.loc[0].values.reshape(1,-1))
```

```
[221]: # house price prediction for new data
y = np.expm1(model.predict(tmp_scaled))
print(y)
```

```
[264503.94660399]
```

```
[222]: new_data['SalePrice'] = y
```

```
[223]: # show the predicted new house:
pd.DataFrame(new_data)
```

```
[223]:
```

	LotArea	YearBuilt	OverallQual	OverallCond	GrLivArea	MiscFeature	\
0	8000	2010	8	7	2000	Shed	
	ExterQual	RoofMatl	SaleCondition	GarageArea	KitchenQual		SalePrice
0	Gd	Compshg	Normal	500	Gd		264503.946604

Random Forest Regression model is performing well with a Mean  $R^2$  of 0.8496, but there might still be room for improvement. Here are some strategies to enhance its performance: 1. Check Feature Importance & remove irrelevant features, but before removing the feature, try running the model with and without certain features, then compare its performance. 2. Experiment with ensemble learning (stacking/blending). 3. Look for external data that can provide more insights into house prices, for example, macroeconomic data (Inflation, interest rates, average salary), more detailed location data (crime rate, school quality, transportation).